

# SKEE: A Lightweight Secure Kernel-level Execution Environment for ARM

APR 11TH, 2016

[论文下载](#)

## Introduction

为内核提供检测与保护  
以前的一些方法：

- Virtualization-based Approaches: TCB 必须足够大，能够处理资源分配、硬件虚拟化。而且虚拟机例如 VMware、XEN 有自身的安全问题。
- Microhypervisor Approaches: 小型的 hypervisor，旨在提供隔离。microhypervisors do not provide a good fit to host security tools that require a relatively large code base.
- sandbox: 需要虚拟化提供隔离，需要 hypervisor 来管理分配 sandbox，因此会增大 TCB 的大小。
- hardware protection: Intel: Software Guard Extensions (SGX). 在 ARM 中没有类似的防护，ARM 提出

TrustZone，但也会增大code base和需要对高权限的Trustzone层的维护

## SKEE

- 完美地与内核隔离
- 提供一个切换到隔离区域的接口
- 能够让安全工具检测内核

## Background

### 32-bit ARMv7

内存转换涉及的三个寄存器

- Translation Table Base Control Register (TTBCR)
- Translation Table Base Register 0 (TTBR0)
- Translation Table Base Register 1 (TTBR1)

TTBCR contains a 3 bit called TTBCR.N

If the value of TTBCR.N is 0, then TTBR1 is not used, otherwise both TTBR0 and TTBR1 are used.

Table I. EFFECT OF TTBCR.N ON ADDRESS TRANSLATION ON 32-BIT ARMv7 USING SHORT DESCRIPTOR FORMAT

TTBCR.N value	Starting address of TTBR1 translation
0b000	TTBR1 not used
0b001	0x8000_0000
0b010	0x4000_0000
0b011	0x2000_0000
0b100	0x1000_0000
0b101	0x0800_0000
0b110	0x0400_0000
0b111	0x0200_0000

### 64-bit ARMv8

- kernel: TTBR1
- user: TTBR0

MSR指令可以使用0寄存器 (XZR)使任何寄存器为0.

Address Space Identifier (ASID): 非全局访问的虚拟内存页，与一个特定的ASID相连。

Memory Management in Virtualization Layer

## **THREAT MODEL, SECURITY GUARANTEES AND ASSUMPTIONS**

### **Threat model**

- 修改、重定位内核执行文件：安全操作
- 通过漏洞修改内核数据或控制流：通过安全工具检测
- 攻击SKEE隔离区域，使防护失效：保证攻击不能到达隔离区域，bypass检测

### **SECURITY GUARANTEES**

- 禁止内核修改内存布局和系统访问权限
- 内核到隔离区域只能通过一个严格控制的switch gate来访问

### **ASSUMPTIONS**

- 整个系统 loaded securely
- 内核支持  $W \oplus X$

## **SKEE DESIGN**

### **SKEE Isolation**

#### **Creating a Protected Address Space**

- 1 去除所有映射到SKEE隔离区域或者内核memory translation tables的entries
- 2 将内核代码和SKEE switch gate映射为只读
- 3 所有其他内存区域（包含内核数据，用户层内存）为PXN

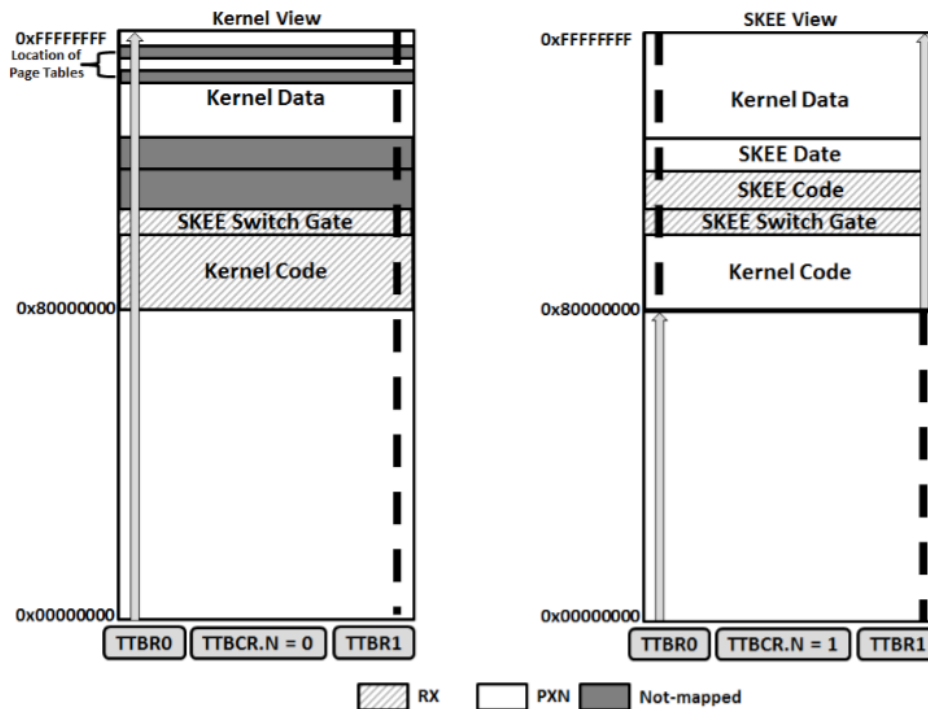


Figure 3. An example of address Space Separation on ARMv7. The vertical arrows show the virtual memory range translated via the corresponding translation table base register. The dashed lines show the memory range that is not addressable via the corresponding translation table base register.

## Restrict Kernel Access to the MMU

SKEE 移除了关于内核的一些控制代码，把它们替换成可以跳到switch gate的hooks

## SKEE Secure Context Switching

switch gate 是 atomic, deterministic and exclusive

- ARMv7

```

1  /* Start of the SKEE Entry Gate */
2  mrs    r0, cpsr                // Read the status register
3  push   {r0}                    // Save the status register value
4  orr    r0, r0, #0x1c0          // Set the mask interrupts bits
5  msr    cpsr, r0                // load the modified value
6
7  mov    r0, #0x11
8  isb                                // Synchronization barrier
9  mcr    p15, 0, r0, c2, c0, 2    // Modify the TTBCR to activate SKEE
10 isb
11
12 mcr    p15, 0, r0, c8, c7, 0    // TLB invalidate
13 isb
14
15 bl     skee_entry              // Jump to SKEE entry point
16 /* End of the SKEE Entry Gate */
17
18 /* Start of the SKEE Exit Gate */
19 mov    r0, #0
20 isb
21 mcr    p15, 0, r0, c2, c0, 2    // Modify the TTBCR to deactivate SKEE
22 isb
23
24 mcr    p15, 0, r0, c8, c7, 0    // TLB invalidate
25 isb
26
27 pop    {r0}                    // Reload status register value
28 msr    cpsr, r0                // Restore the original status register
29
30 bl     kernel_entry            // Jump back to the kernel
31 /* End of the SKEE Exit Gate */

```

Figure 4. SKEE Switch Gate on ARMv7

- ARMv8

```

1  /* Start of the SKEE Entry Gate */
2  mrs    x0, DAIF                // Read interrupt mask bits
3  str    x0, [sp, #-8]!          // Save interrupt mask bits
4  msr    DAIFset, 0x3            // Mask all interrupts
5
6  mrs    x0, ttbr1_el1           // Read existing TTBR1 value
7  str    x0, [sp, #-8]!          // Save existing TTBR1 value
8
9  msr    ttbr1_el1, xzr          // Load the value Zero to TTBR1
10 isb
11
12 tlbi    vmalle1                // Invalidate the TLB
13 isb
14
15 adr    x0, skee_entry           // Jump to SKEE entry point
16 br    x0
17 /* End of the SKEE Entry Gate */

```

Figure 5. SKEE Entry Gate on ARMv8

```

1  /* Start of the SKEE Exit Gate */
2  nop                                //no operation
3  nop                                // Fill the page with no operations to
4  nop                                // align the last instruction with the
5  nop                                // bottom of the isolated page boundary
6
7  msr    DAIFset, 0x3            // Mask all interrupts
8
9  ldr    x0, [sp, #8]!           // Reload kernel TTBR1 value
10 dsb    sy
11 msr    ttbr1_el1, x0           // Restore TTBR1 to kernel value
12
13 /*-----Isolated Page Boundry-----*/
14
15 isb
16 tlbi    vmalle1                // Invalidate the TLB
17 isb
18
19
20 ldr    x0, [sp, #8]!           // Reload interrupts mask bits
21 msr    DAIF, x0                // Restore interrupts mask bits register
22
23 ret
24 /* End of the SKEE Exit Gate */

```

Figure 6. SKEE Exit Gate on ARMv8

- Using ASID for Faster Context Switching

```

1  /* Start of the SKEE Entry Gate */
2  push    {lr}                // save the return address
3  mrs     r0, cpsr             // read the status register
4  push    {r0}                // save the status register value
5  orr     r0, r0, #0x1c0       // set the mask interrupts bits
6  msr     cpsr, r0             // load the modified value
7
8  mrc     p15, 0, r0, c13, c0, 1 // read current CONTEXTIDR
9  push    {r0}                // save the CONTEXTIDR value
10 mov     r0, #0
11 mcr     p15, 0, r0, c13, c0, 1 // set CONTEXTIDR to 0
12 isb
13
14 mov     r0, #0x11
15 isb                          // synchronization barrier
16 mcr     p15, 0, r0, c2, c0, 2 // modify the TTBCR to activate SKEE
17 isb
18
19 L1:
20 mrc     p15, 0, r0, c13, c0, 1 // read current CONTEXTIDR
21 cmp     r0, #0               // compare current CONTEXTIDR with 0
22 bne     L1                   // loop back if CONTEXTIDR is not 0
23
24 bl      skee_entry           // jump to SKEE entry point
25 /* End of the SKEE Entry Gate */
26
27 /* Start of the SKEE Exit Gate */
28 mov     r0, #0
29 isb
30 mcr     p15, 0, r0, c2, c0, 2 // modify the TTBCR to deactivate SKEE
31 isb
32
33 L2:
34 mrc     p15, 0, r0, c2, c0, 2 // read current TTBCR
35 cmp     r0, #0               // compare current TTBCR with 0
36 bne     L2                   // loop back if TTBCR is not 0
37
38 pop     {r0}                 // reload CONTEXTIDR value
39 mcr     p15, 0, r0, c13, c0, 1 // write original CONTEXTIDR
40 isb
41
42 L3:
43 mrc     p15, 0, r0, c13, c0, 1 // read current CONTEXTIDR
44 cmp     r0, #0               // compare current CONTEXTIDR with 0
45 beq     L3                   // loop back if CONTEXTIDR is 0
46
47 pop     {r0}                 // reload status register value
48 msr     cpsr, r0             // restore the original status register
49
50 pop     {lr}                 // reload the return address
51 movs    pc, lr               // jump back to the kernel
52 /* End of the SKEE Exit Gate */

```

Figure 7. A Faster SKEE Switch Gate on ARMv7

## **Kernel Monitoring and Protection**

- 可以 trap kernel critical events
- 可以访问内核内存
- 可以进行内核内存保护

## **Security Analysis**

- 32-bit ARMv7: Samsung Note4
- 64-bit ARMv8: Samsung Galaxy S6、 Samsung Galaxy Note5

## **IMPLEMENTATION AND evaluation**



Table III. SKEE BENCHMARK SCORES ON ARMv7

Benchmark	Original	SKEE	Degradation (%)
CF-Bench	30933	29035	6.14%
Smartbench 2012	5061	5002	1.17%
Linpack	718	739	-2.93%
Quadrant	12893	12552	2.65%
Antutu v5.7	35576	34761	2.29%
Vellamo			
Browser	2465	2500	-1.42%
Metal	1077	1071	0.56%
Geekbench			
Single Core	1083	966	10.8%
Multi Core	3281	2747	16.28%

Table IV. SKEE BENCHMARK SCORES ON ARMv8

Benchmark	Original	SKEE	Degradation(%)
CF-Bench	75641	66741	11.77%
Smartbench 2012	14030	13377	4.65%
Linpack	1904	1874	1.58%
Quadrant	36891	35595	3.51%
Antutu v5.7	66193	67223	-1.56%
Vellamo			
Browser	3690	3141	14.88%
Metal	2650	2540	4.15%
Geekbench			
Single Core	1453	1235	15.00%
Multi Core	4585	4288	6.48%

Table V. SKEE APP LOAD DELAY ON ARMv7

App	Original	SKEE	Overhead (%)
Temple Run 2	9.31	10.33	10.96%
Hill Climb Racing	3.66	3.71	1.37%
Angry Birds	4.72	4.79	1.48%
Crossy Road	4.81	5.24	8.94%
Subway Surf	5.45	5.95	9.17%

Table VI. SKEE APP LOAD DELAY ON ARMv8

App	Original	SKEE	Overhead(%)
Temple Run 2	6.08	6.58	8.22%
Hill Climb Racing	2.42	2.73	12.81%
Angry Birds	4.12	4.32	4.85%
Crossy Road	3.42	3.80	11.11%
Subway Surf	4.42	4.71	6.34%

Table VII. SKEE ADDED SECURITY CHECKS BENCHMARK SCORES

Benchmark	Original	SKEE	Degrad. (%)	SKEE + Sec. Checks	Degrad. (%)
CF-Bench	63273	58903	6.91%	57250	2.81%
Smartbench	15820	15217	3.81%	15104	0.74%
Linpack	1849	1697	8.22%	1560	8.07%
Quadrant	31429	30843	1.86%	29330	4.91%
Antutu v5.7	65242	62866	3.64%	58658	6.69%
Vellamo					
Browser	4659	4256	8.65%	4350	-2.21%
Metal	2158	2139	0.88%	2081	2.71%
Geekbench					
Single Core	1508	1342	11.00%	1340	0.15%
Multi Core	4566	4388	3.90%	4207	4.12%